

Chapter 3: Identifying, Acquiring, and Representing Text Quantitatively

Justin Grimmer* Margaret E. Roberts[†] Brandon Stewart[‡]

January 8, 2018

Abstract

Reading notes: Thanks for reading this chapter! A few things to flag as you go through: I'm not sure that having the examples at the end is the right ordering – the examples perhaps should be interspersed in the text. I'd be interested to get your thoughts on that. I'm also missing a few things I think we should probably put in – dependency parsing, optical character recognition, plagiarism detection and other measures of similarity. Other things that you think we might be missing would be helpful. The collecting data and representing it quantitatively seems disjointed – I'd be interested to know ways those could better fit together. Last, I'm interested in whether certain sections need more/less detail – most of these sections you could spend a book on, so I'm wondering what we need to include and what we should exclude.

1 Introduction

Once you start to see text as a source of data, you quickly start to realize that the opportunities for studying the social world through the lens of text abound. For economists, text data present themselves in company earning reports, consumer complaints, and product reviews. For political scientists, text data appear in speeches, press releases, and debates. For historians, text data are buried in archives and interviews. Even outside of social science, text data can facilitate learning in medicine and public health through medical records, manufacturing through accident reports, and the humanities through the digitization of literature.

*Associate Professor, Department of Political Science, Stanford University

[†]Assistant Professor, Department of Political Science, University of California at San Diego

[‡]Assistant Professor, Department of Sociology, Princeton University

But how can we use this abundance of text data and burgeoning tools that facilitate the automated analysis of text to make inferences about the world? Collecting some text and doing some analysis is straight forward, but actually learning something about the human processes that created the text is more complicated. As we discussed in Chapter 2, the collection and analysis of text should be problem-oriented. In this chapter, we focus on how to set up the analysis of text with the intention of shedding light on a social science question. We begin by discussing how a principled way of gathering and sampling text is imperative for making inferences. After gathering text, we explain how text can be represented with numbers, and how the scientific question of interest is important in making decisions about how to represent text. We end with some examples of how different authors have made decisions to represent text differently in order to answer their own research question.

2 Sampling Biases and Selecting Corpora

Once the researcher has identified a domain to explore or a specific research question to test, the first step toward answering this question is to select documents to use as data. Because the population of interest will depend on the research question, sampling bias cannot be defined across corpora, but are specific to each individual research question. For inferences to generalize to a broader population of interest, the texts ideally will be representative of the individuals, institutions, or social phenomena of interest to the researcher. However, even a sample with known biases can sometimes be useful for generating social science insights, as long as the biases are explored and documented.

In this section, we provide a step-by-step approach to thinking through common sources of bias in text corpora. First, researchers should think carefully about the broader population of interest and the quantities they want to measure. In selecting a corpus, researchers should consider reasons documents or text might be missing from the corpus, including the

physical and incentive-based constraints on those providing and writing the documents. Last, researchers could consider how their method of sampling may bias their data, particularly if they relying on third-party software or are using their own statistical techniques or keywords to select subsets of a larger corpus.

2.1 What are the populations and quantities of interest?

Researchers in the social sciences ask questions that focus on wide variety of populations, domains, and time periods. They should think carefully about what social phenomena they hope to reflect and what aspects of the phenomena they hope to measure when selecting a corpus. For some research questions, text data may be used similarly to population surveys, where researchers hope to reflect the underlying opinions, activities, or demographics of a citizen population. For example, scholars have begun to use social media data like Twitter to describe everything from political opinions (O'Connor et al., 2010) to movie preferences (Asur and Huberman, 2010), the spread of the flu (Lampos and Cristianini, 2010), and personal happiness in a population (Golder and Macy, 2011; Dodds et al., 2011). To accurately gauge underlying quantities of interest in the broader (both online and offline) population, these analyses depend on accurately accounting for the ways in which online populations are different from offline populations. Even within online populations, those who choose to engage in discussions about politics may be very different than those choosing to engage in discussions about entertainment, all which must be taken into account when generalizing to the population that does not discuss the topic within the data (Barberá and Rivero, 2014).

In other cases, the social phenomenon of interest are the texts themselves. For example, Blaydes, Grimmer and McQueen (2014) explore the divergence in political thought between Middle Eastern and Christian societies. To reflect both Muslim and Christian political thinkers, they gather political advice texts from the 6th to the 17th centuries from both the Islamic tradition and Christian Europe. The authors' goal is not to gather *all* Muslim and

Christian writings, or to reflect all thoughts on these subjects within these populations, but rather to reflect those with the most influence on political thought in this time period. In the digital humanities, analyses may be focused on exploring thematic trends in a particular genre, for example, Danish ghost stories (Abello, Broadwell and Tangherlini, 2012). The selection of texts in these cases will often rely on the ability to distinguish between stories that contain a subject like ghosts from those that do not.

At times, corpora that are known to be starkly biased can still provide a window into a social processes. For example Gill and Spirling (2015) use historical records and leaked data to understand what types of information the United States government keeps secret and the process of declassification. Of course, the authors do not have access to *all* classified information, they each only see that which has been leaked or that which has eventually been declassified. But in comparing information that is unclassified to that which was classified at some point, the authors can infer the types of topics that remain secret and can offer insight to what might be missing.

Almost no text corpus is fully representative of the social phenomenon of interest, and in each of these examples the authors examine and document the potential biases that may remain. In the subsequent sections, we address four types of sampling bias that could arise in text corpus, drawing examples from each of these applications. These types of sampling bias are not meant to be exhaustive, but provide a starting checklist of things to consider for researchers considering a new text corpus.

2.2 Resource Bias

Texts are expensive to produce, gather, and collate and every individual, institution, or group in the population of interest may not have the resources to record, preserve, and disseminate texts. Researchers should consider which portions of their population of interest have the capabilities to write and disseminate texts so that researchers can access them. These issues

are especially salient when the population of interest expands beyond the texts themselves; for example, when researchers use text data produced by individuals to make inferences about a broader population of interest. Portions of the population may be illiterate and these people will be completely unrepresented in such texts. Texts downloaded from the Internet may underrepresent those who do not have access to the Internet, and some of the poorest areas in even the most developed countries are much less connected to the Internet and will be underrepresented in social media studies (Norris, 2001).

Resources may also affect the availability of historical data. Historical events may be more likely to be recorded if they occur in areas where the press is present (Snyder and Kelly, 1977; Danzger, 1975). Archives can sometimes be a combination of documents that citizens found easy to store and may be woefully unrepresentative of the whole set. Texts may have been lost, burned, or not stored, and stories may have never been transcribed. Communities and individuals with the capabilities to store and preserve texts will likely be very different than those without, and as such historical documents should not be considered a random sample of a larger set.

Government document availability may also reflect local resources. Some governments may have the personnel to transcribe meetings and make them available to the public, others will not. Even governments in the U.S. exhibit high variability in their capabilities to store and make available documents. For example, local police reports are sometimes handwritten, making them difficult to make available to researchers, while others are entered directly into a computer. The ways in which government documents are stored can make them more costly for researchers requesting them through Freedom of Information Act laws, as fees are charged to researchers for identifying and duplicating the documents of interest.

2.3 Incentive Bias

In some cases, not only resources but human incentives will affect the availability of documents and may explain corpus missingness. Individuals and organizations have incentives to fail to record, hide or destroy evidence that could cast them in a negative light. Individuals may be more likely to post social media that reflects the happiest or most successful aspects of their lives. For researchers studying interactions on social media, this selection bias may not be an issue, but for those who are trying to measure emotional states of users, such selection may bias results. Similarly, politicians and governments may force removal of news and social media that undermines them, censoring others or self-censoring to frame political conversations (House, 2015; Boydston, 1991; Gup, 2008; King, Pan and Roberts, 2013; MacKinnon, 2008; Pomerantsev, 2014).

While transparency laws and initiatives may make texts more accessible to researchers, it also changes the incentives of those whom it affects. E-mail transparency within organizations may create incentives for members of these organizations to talk on the telephone about sensitive issues or hold informal meeting where discussions are not recorded. Transparency initiatives might make government leaders curtail their political participation they are afraid of making missteps (Malesky, Schuler and Tran, 2012). When possible, researchers should conduct extensive interviews with those writing, storing, and making documents available to understand the incentives behind document collection and the process through which the data are made available.

2.4 Medium Bias

The technologies and types of mediums in which texts are recorded will play an important role in the types of content that will be reflected in them. Twitter, for example, only allows 140 characters in one tweet, necessitating users to use abbreviation and links for further treatment

of the topic. However, 140 characters has very different implications across languages – in denser languages like Chinese, users will be able to express more content in one post than in languages that necessitate many characters for each word like English. Further, the mode of expression in social media change with advances in technologies. As users have been able to use pictures, videos, and live streams online, the requirements of the accompanying text changes. Users interested in text analysis in social media should be aware of other types of metadata – links, pictures, and videos – when describing their content and how technological capabilities have changed over time.

The technology and evolution of the medium can create different cultures for users so that the text can only be understood in the context of the platform. Snapchat, a social media platform where posts are deleted immediately after they are read, will create different incentives for users than Twitter, where posts are public and more permanent. As most online platforms present posts to users that depend on their past history or their predicted characteristics, even within the same medium, users might have very different perspectives even about what is being discussed on the platform. Politicians may try to drive steer the conversation toward topics that reflect better on them (Franco, Grimmer and Lee, 2016). Researchers should consider biases that might be produced from individual users' different experiences within the same platform and consider how the topics discussed might be strategically manipulated.

Text outside of social media is similarly influenced by its medium. Meeting transcripts or notes may reflect the content of the conversation, but not tone of voice, body language, or other forms of expressed emotion. Handwritten notes may contain drawings or doodles, computer written messages may contain emojis. When possible, researchers should read and interact with texts in their original contexts to understand how the social events of interest are translated by the medium into text, even if text analysis tools only take into account the text themselves.

2.5 Algorithmic Bias

As texts are sometimes selected using statistical methods, algorithmic bias can sometimes affect the selection of corpora. Say, for example, a researcher is interested in doing a quantitative analysis of ghost stories (Abello, Broadwell and Tangherlini, 2012). To do so, they might select every book in the library that includes in the title “ghost.” However, this selection might be a biased sample of ghost stories. Some book titles that contain the word ghost may have nothing to do with ghost stories; for example, *Hungry Ghosts: Mao’s Secret Famine* deals with the famine in China rather than with ghost stories themselves. Similarly, many books that don’t contain the word “ghost” may still be ghost stories, *A Christmas Carol* is a ghost story without “ghost” in the title.

While these errors may at first seem random, they are often systematic. “Ghost” may orient the analysis toward one type of genre, missing books about phantoms which may have different themes. The word ghost might miss books about ghosts in other languages.

This type of keyword selection of corpuses is a frequent problem in the analysis of text because researchers often want to look at a subset of a larger corpus. The problem is made worse by the fact that humans often have difficulty thinking of words off of the top of their head, they have limited recall, so may miss some of the most important words on which to select. Researchers might want to use methods that use the text itself to generate candidate keywords, or supervised learning methods that can identify texts of interest (King, Lam and Roberts, 2014; D’Orazio et al., 2014). Regardless, researchers should be transparent about how they selected the texts and what potential biases this might produce.

While keywords are a simple algorithm for selecting texts, more complicated algorithms are also used for selecting corpuses. Further, algorithmic bias is broader than a researcher’s own method of selection as researchers often rely on third parties to provide data for them. Researchers who access data through Application Programming Interfaces (API) or other search functions should be as aware as possible of the underlying process that retrieves search

results from queries. For example, some interfaces may account ignore capitalization, others might be sensitive to it. Some may provide a random selection of documents, others may return the most popular. The process through which the API makes documents available will affect the population the text represent and the research questions the text can answer.

2.6 Time-series problems with text selection

The biases in the selection of texts may change over time, making it difficult for researchers to discern whether shifts in the texts over time are created by changing opinions or styles of individuals within the population or by shifts in the sample. For example, when Facebook first started, it was originally open only to particular colleges, but then was expanded to the public. Without considering the sample, a researcher might come to the conclusion that people have become less interested in education over the years. However, this change would be due to the changing sample of . Policy shifts in data transparency can create important discontinuities in government document samples that may appear to the researcher as if the substance of governance is shifting even if the underlying documents are static over time. Researchers should be careful not only to consider static bias but also how biases in document samples could vary over time to account for demographic, resource, medium, or incentive drift.

3 Representing Texts

After you have identified a corpus of interest, the next step in analyzing the text is to retain only the necessary information from each document for your analysis. Texts are complex – they encode meaning through word choice, phrasing, and the organization of concepts. There are many different features of text that an analyst could encode as data. However, not all of the complexity of text is necessary for any particular analysis. Because quantitative analyses

of text are often focused on characterizing a few general characteristics of a large document set – characterizing the haystack instead of a straw of hay – our analyses will not require us to retain all information from the texts. Reducing the complexity of the text will be a key step in allowing the analysis to focus on only the aspects of the text that are important to the task at hand.

Here we will provide you with a recipe of commonly-used tools for simplifying texts to represent them quantitatively. We stress that there is not a one-size fits all solution to preprocessing texts. The steps you take to represent the text will necessarily depend on your specific research question and focus. Indeed, the choices you make about how to represent text will affect your subsequent analyses, so these choices should be made deliberately and with care (Denny and Spirling, 2016). We will identify some points you might want to consider while deciding how to preprocess the texts before your analysis.

Before you preprocess the text, the first decision you must make is what is the *unit* for your analysis, what is the basic unit of text that you want to describe? In many cases, the unit of analysis is the document – most researchers want to categorize documents from newspaper articles to social media posts to journal articles. However, in some cases the unit of interest might be smaller or larger than the document – the analyst may want to break apart the document in order to categorize sections, paragraphs, or sentences or may want to describe a collection of documents such as journals, a newspaper, or a book. Whatever text you decide to use as your basic unit of analysis, the collection of these units is called a *corpus*, or the full dataset on which you will run your analyses.

The purpose of the analysis is to describe the units of text, for example, to put each into specific categories or to summarize the main themes of all of the units in the corpus. To do this quantitatively, each text must be turned into data, described by their *features*. The features of the text are any aspect that can be measured on all of the texts. Features are commonly described in data analysis as variables. One simple example of a feature in

texts are the number of a particular word; for example a count of how many times each text unit contains the word “political.” In addition, features can be metadata about the text – the time period the text was written in, the author of the texts, or the newspaper the text was printed in for example. Features can also be a derivation from the words; for example, whether or not the text contains a proper name or in other cases a low-dimensional projection of the words. We will describe each of these types of features in this chapter in turn.

3.1 Tokenization

The basic problem of representing the text is deciding which features to use to describe your text units in your analysis. Texts are often thought of as *high-dimensional* in that there are many possible features of any given text. Texts not only contain information about what words are in them, but also the order in which these words appear, and other metadata associated with the document. Typically, there is too much information within the text to be represented efficiently by the data.

First, the researchers have to decide which words should be included as individual features, a process commonly known as *tokenization*. Many methods for text analysis use what is known as the “bag of words” assumption, the idea that features should include counts of unique words in the corpus, often called “unigrams,” which are separated by spaces. With the bag of words assumption, each individual word is a feature or token, and the number of each word is used to describe the text, but nothing is retained about word order.

The “bag of words” assumption is a useful starting place for tokenization, but is not ideal for all contexts. In many languages, for example Chinese, Japanese, and Lao, words are not separated by spaces; instead, words are inferred by the reader from the context within the sentence. In these languages, the author of the document has not provided the tokenization, so the analyst must decide how the text should be segmented out into features. For each

of these language, software has been developed to tokenize the words, or separate them by spaces. The bag of words methods can then be applied directly.

In other cases, “unigrams” or individual words discard important word order that might be important to retain. Some concepts bridge words, a common example is “White House,” where the meaning is lost when the phrase is tokenized into “white” and “house.” For some applications, word order may be particularly useful in predicting the categories of the documents or describing the main themes of the corpus. In this case, the analyst can retain “n-grams,” or tokens that include n words. For example, the researcher may retain all “bigrams,” or consecutive sets of two words, or “tri-grams,” consecutive sets of three words. In this case, each unique bigram or trigram is a feature, and the number of each bigram or trigram is used to describe each unit of text.

The researcher can include n-grams in two ways. First, she could include all consecutive sets of two words that appear in the corpus, in addition to or instead of unigrams. This might be most useful in cases where she performs another step of automated feature selection, where the computer decides what of a given set of features to retain, something we will discuss later in the book. Alternatively, the researcher might find it useful to retain only certain bigrams and trigrams that she anticipates will be useful to her analysis. For example, she might know that “White House” or “International Relations” often appear together and therefore may specify which bigrams should be included as one token. A variety of approaches exist to suggest bigrams and trigrams that frequently appear together for inclusion within a corpus (Handler et al., 2016).

For now, think of tokenization as a process that either inserts spaces in between words (such as in the case of Chinese, Japanese or Lao texts) or removes spaces in between words to incorporate n-grams. Each space indicates that information about the word before or after that space is completely discarded, the only information retained are the words within the spaces. Below, we show an example of how a document might be tokenized in a Chinese

document and in an English document where the researcher chooses to include select bigrams.

Document 1

“ 人生就像蒲公英，看似自由，却身不由己。有些事，不是不在意，而是在意了又能怎样。”

“ 人生 就 像 蒲 公 英 ， 看 似 自 由 ， 却 身 不 由 己 。 有 些 事 ， 不 是 不 在 意 ， 而 是 在 意 了 又 能 怎 样 。 ”

Document 2

“Climate change constitutes a serious threat to global security, an immediate risk to our national security ... And so we need to act – and we need to act now. –President Obama”

“Climatechange constitutes a serious threat to global security, an immediate risk to our national security ... And so we need to act – and we need to act now. –President Obama”

3.2 Discarding Complexity in Word Features

Once the words in the document have been tokenized, analysts often discard features that are unimportant to the analysis. For example, if the researcher is not interested in how many commas, apostrophes, or periods appear within the text, then choosing to delete all punctuation might be useful since it will decrease the total number of features in the data. Second, the analyst may decide to remove *stopwords* or common words used across documents that do not give much information about the categories the researcher wants to place the texts in. In English, common words such as “the,” “that,” and “and” might be

removed from the documents to reduce the size and complexity of the text. All languages have such commonly used words that function as prepositions or articles. Lists of stopwords in many languages are available in packages that provide text analysis tools, and statistical software programs for text analysis often have automated functions to remove punctuation and stopwords.

The researcher may also want to get rid of formatting cues that often exists within data from other sources. For example, section headings, page numbers and html tags, may not be important to the purpose of the research, and, if so, should be discarded before the analysis.

Researchers may combine some of the features that describe the text by removing the endings of words, or reducing words to their root. Those working with texts in the English language typically discard the ending of conjugated verbs and plural nouns to merge similar words into one root, in a process called “stemming.” For example, “family,” “families,” “family’s” and “familiar” would all be replaced with “famili” in the text. This stemming process reduces the number of features that the analyst has to deal with, while not discarding significant information since words that begin with famili all have similar meanings. There are many stemming algorithms available in a wide range of software; among the most common is the *Porter* stemming algorithm (Porter, 1980).

A more complex version of stemming is lemmatization, which rather than simply removing word endings, reduces words to their “lemma,” or base root. For example, “see,” “saw,” and “seeing” would not be reduced to one simplified word with stemming, but with lemmatization these words will all reduce to the word “see.” Lemmatization is typically slower than stemming because the algorithm considers the word in its context (Grimmer and Stewart, 2013; Jursfsky and Martin, 2009; Manning, Raghavan and Schütze, 2008). However, in languages such as German, where conjugations are not always represented by adding an ending to a word, lemmatization can be more useful than stemming.

Below we provide Document 2 from above, but where the text has been converted to

lowercase, tokenized, removed of punctuation and stopwords, and stemmed, in that order. As of yet, the researcher has already made important decisions about which features he will ultimately include within the analysis. Note that the order in which these steps are taken matters as well. For example, if stopwords are removed before tokenization, the analyst will end up with different bigrams and trigrams than if they are removed before. Researchers should examine the texts after pre-processing to ensure that the result of pre-processing will make sense for their analysis.

When pre-processing steps the research wants to do are not available in software, researchers can use *regular expressions* to manipulate text to customize preprocessing to their own dataset. Regular expressions may also be useful for discarding unnecessary text such as formatting or html that interrupts the text and is unnecessary to the analysis. Regular expressions are a language for searching general string patterns. They form the basis for string search for many programming languages, search engines, and text editors. They also provide functionality for replacing or removing unwanted text. For more information on how to use regular expressions, along with examples of regular expressions see Jursfsky and Martin (2009, Chapter 2).

Document 2

“Climate change constitutes a serious threat to global security, an immediate risk to our national security ... And so we need to act – and we need to act now. –President Obama”

“Climatechange constitutes a serious threat to global security, an immediate risk to our national security ... And so we need to act – and we need to act now. –President Obama”

“climatechang constitut serious threat global secur immedi risk national secur need

act need act now presid obama”

3.3 The Document-Feature Matrix

Once you have tokenized the corpus, removed unnecessary words, word endings, and formatting, you are ready to turn text into numeric data. The most basic way to do this is first to identify all unique tokens within your dataset. We will call this the “vocabulary.” Each unique token becomes a variable within your dataset, and the the number of times each unique token appears within the text becomes your data. In other words, for each unit of text, you will count how many times each token is used within that text and record that as your data.

The result of this is what we call a document-feature matrix. In this matrix, there are as many rows as there are unique units of text, and as many columns as there are features, or unique tokens. Each element of this matrix is the number of times that unique token appears in each unit of text. In general, since most documents only contain a few of the many unique tokens within a corpus, the document-feature matrix is relatively *sparse*, or contains lots of zeros.

For example, say that we have the three documents in a corpus of Presidential quotes, below:

Presidential Quotes Corpus

“Climate change constitutes a serious threat to global security, an immediate risk to our national security ... And so we need to act – and we need to act now. –President Obama”

“Whenever America fights for the security of our country, we also fight for the values of our country.–President Bush”

“Climate change is more remote than terror but a more profound threat to the future of the children and the grandchildren and the great grandchildren – President Clinton”

We could then put them through the same preprocessing routine as we did with Document 2, above. This would create three documents, each having been converted to lowercase, punctuation removed, stopwords removed, stemmed, and converted “climate change” into a n-gram “climatechang”. After doing this, we would document the number of unique words within the corpus, the *vocabulary*. Overall, after being preprocessed these three documents have 29 unique words. The document-feature matrix would therefore consist of three rows (one for each document) and 29 columns (one for each unique word). The entries of those rows would count the number of times each word appeared in each document, see below. See the document-feature matrix for the Presidential quotes corpus below.

	clinton	great	grandchildren	children	futur	profound	terror	remot	bush	valu	also	countri	fight	america	whenev	obama	presid	now	act	need	nation	risk	immedi	secur	global	threat	serious	constitut	climatechang
text1	1	1	1	1	1	2	1	1	1	2	2	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
text2	0	0	0	0	0	1	0	0	0	0	0	0	1	0	1	1	2	2	1	1	1	0	0	0	0	0	0	0	
text3	1	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	1	1	1	

Table 1: Document-feature matrix from the Presidential quotes corpus

Once we have the document-feature matrix, there are many operations we can perform on it that can help summarize the text. To take two simple examples to start, if we sum each of the columns, we will have a summary of how many of each unique token appear within the corpus. This could then be inputted into a word cloud or a table to understand the most common words within the corpus.

Or, we could sum each of the rows of the document-feature matrix, which would tell

clinton	1
great	2
grandchildren	1
children	1
futur	1
profound	1
terror	1
remot	1
bush	1
valu	1
also	1
countri	2
fight	2
america	1
whenev	1
obama	1
presid	3
now	1
act	2
need	2
nation	1
risk	1
immedi	3
secur	1
global	2
threat	1
serious	1
constitut	1
climatechang	1

Table 2: Result of summing the columns of the document feature matrix.

text1	17
text2	11
text3	12

Table 3: Result of summing the rows of the document feature matrix.

us how many unique tokens each document has. The maximum of this sum would be the length of the longest document, and the minimum of this sum would be the length of the shortest document. While these are basic examples, much of the rest of this book describes algorithms that perform operations on this document-feature matrix.

3.4 Adding features

Once we have the document-feature matrix, we can add additional features outside of unique tokens within the text, creating additional columns of our matrix that provide information about the documents within the corpus. One example is that we can add metadata to our document-feature matrix. The documents not only have words that describe them, but have other information about the context in which they were written. For example, the date or time period it was written, the book or newspaper the article appeared in, the author of the document, or other attributes of the author such as gender, age, or political party, might be important information to add to the feature set. These attributes can be used in conjunction with the text to classify the documents, or might be used as primary inputs in the analysis.

There might be other, non-textual features that play an important role in communication that we might want to add to the corpus. For example, we might have the text of the words

within a news broadcast, but want to add information about the tone of the speaker, or the music playing in the background (Schonhardt-Bailey, N.d.). Or, we might have the transcript of a meeting that includes non-verbal cues such as laughter, pauses, or applause. If non-verbal, audio, or visual cues are important to our analysis, they should be added as features in our data.

3.5 Leveraging text complexity

So far, we have described very basic processes of turning text into data: first tokenizing the document using n-grams, next removing unnecessary complexity, and creating a document-term matrix. For the most part, our assumption so far has been that the order and context of the words is not important – we have been using the “bag of words” assumption to simplify the text.

While as we will discuss later, this simple approach is surprisingly effective for many research questions, it may be too simple for some tasks. One of the problems with the bag of words approach is that words can be ambiguous and imprecise, causing confusion because words are taken out of context. For example, one feature of the document-feature matrix might be “crack.” Does this refer to a break, as in a crack in a window pane, “cracking” or solving a problem, or the nickname for cocaine? Two words might also have the same meaning, but have two separate features. For example, “Dairy Queen” could also be called “DQ” and similarly “McDonald’s” “Mickey D’s.” While these difference could be unimportant for some problems, their separation might be problematic for others.

In this section, we will describe how to add information about the order and word context to the document-feature matrix, sometimes called “resolving” ambiguities within the data. Their resolution can be used to add information to the document-feature matrix by adding information about context, or simplifying the document-feature matrix by combining two separate features that should be considered one feature.

3.5.1 Parts of Speech Tagging

Tagging the parts of speech within a sentence can be useful for resolving ambiguities in the data, either by combining words that have the same parts of speech, or disambiguating words that are the same by have different meanings depending on their parts of speech. Parts of speech tagging assigns a part of speech to each word within a document. For example, crack could be tagged with “verb” in “I cracked with window” or a noun in “The window has a crack.” Tagging of each word is either done by a rule-based algorithm, but more often is done by machine learning algorithms that use an already tagged corpus to learn tags for each word based on their relationship to each other (Jursfsky and Martin, 2009, pg 297). For example, the Penn Treebank is a corpus of 4.8 million words from a variety of English language sources tagged with 45 different parts of speech tags and is often used as a training set for parts of speech tagging in English (Marcus, Marcinkiewicz and Santorini, 1993)

As a simple example, we tag the parts of speech of the quote from President Bush from above, shown below. The parts of speech tagger finds four words as nouns – “security,” “country,” “values,” and “country” – and three words as proper nouns – “America,” “President,” and “Bush.” Note that because of the context of the document, the parts of speech tagger is able to distinguish from the plant “bush” from the last name “Bush.”

Whenever ADV America PROPON fights VERB for ADP the DET security NOUN of ADP our ADJ country NOUN , PUNCT we PRON also ADV fight VERB for ADP the DET values NOUN of ADP our ADJ country NOUN .- PUNCT President PROPON Bush PROPON
--

Parts of speech tagging can be useful in augmenting the feature set. For example, while the bag of words approach would merge both the noun and verb versions of “crack” into one feature, after identifying the parts of speech, the research could include one feature

for crack-verb and one feature for crack-noun. Researchers may also be interested in only using particular types of speech for an analysis or components of an analysis. For example, O'Connor, Stewart and Smith (2013) develop a model to characterize the topics of events that happen between a particular object and subject within a sentence. For this, they want to characterize the verb in between a selection of objects and subjects in the text. In this case, parts-of-speech tagging allowed them to only input verbs into a portion of the topic model.

3.5.2 WordNet

Parts of speech tagging facilitates larger groupings of words into groups or common themes. For example, researchers have developed large databases that document relationships between words that have the same part of speech in the English language to facilitate linking similar words together. The largest of such English language project is WordNet (Miller et al., 1990). WordNet allows researchers to query a database of nouns, verbs, or adjectives to retrieve information about a given word. The information returned includes definitions and synonyms as well as relationships with other words; for example, allowing researchers to extract more general concepts that subsume a word in the text and others (for some good examples, see Jursfsky and Martin (2009, pg 603)).

WordNet can allow for the simplification of the document-feature matrix by identifying hierarchies of words. For example, the researcher might only be interested in whether a word indicated that the noun was a car, and therefore would want convertible, truck, and car to be merged into one feature “car.” Some algorithms focused on measuring the readability of text use WordNet to identify the number of causal verbs and causal predicates, an indicator of how complicated the text is (Graesser et al., 2004). Other researchers have used WordNet to help human coders identify a set of emotional words in text, finding synonyms for emotions such as anger, disgust, fear, or joy (Strapparava and Mihalcea, 2008).

3.5.3 Named-Entity Recognition

Last, parts of speech tagging can facilitate the recognition of people and places within text. Using similar methods of classification from tagged texts as parts of speech tagging, algorithms have been developed to tag proper nouns within text (Manning et al., 2014). Named-entity recognition (NER) can also be used to simplify or complicate the document-feature matrix. For example, two people with the same name might be distinguished by where they live or the places that they are associated with. Anne Hathaway may mean a different person in the context of William Shakespeare (wife of William Shakespeare), than in the context of *The Devil Wears Prada* (American actress). Other names mentioned within the document could therefore be used to separate two individuals in different contexts. NER could also be used to simplify the document-feature matrix. If an analyst is only interested in people mentioned within the documents in order to form a network of co-occurrences, after running NER they could subset the vocabulary to only terms that are recognized as proper nouns. The network of co-occurrences is then defined as the $D^T D$, where D is the document-feature matrix.

For a simple, example of named-entity recognition, we used NER from the Spacy Python package to identify named entities within the quote from President Bush. Note that not only does the algorithm recognize the entities within the quote, it also identifies the type of entity – in this case that “America” is a geo-political entity and that “Bush” is a person. It also records the location of the entity in number of characters from the beginning of the document.

(‘America’, 9, 16, ‘GPE’)
(‘Bush’, 113, 117, ‘PERSON’)

3.6 Vector representations of documents and similarity metrics

Once the document-feature matrix has been created, information about the documents and their relationships to one another become easy to extract. Each document is now represented by a vector of numbers, which signify the number of each of the features it includes. Mathematical operations are now easy to perform on this newly-constructed matrix with documents represented as vectors.

Most simply, similarity metrics can be computed on the document set to see how similar any pair of document is to the other. One way to do this would be compute the inner product between two documents. This would involve multiplying each corresponding entry in two document vectors together and then summing these terms together. The more non-zero terms in the two vectors overlap, the larger the resulting number produced is, indicating higher levels of similarity.

One problem with the inner product is that it is length dependent. If I duplicated all words so they were written twice as frequently in one document, my inner product would increase even though the similarity of the two documents has arguably remained unchanged. More frequently, then researchers use the cosine of the angle between the two documents as a measure of similarity, called *cosine similarity*.

Cosine similarity is more simple than it sounds. Imagine you had only two words, as in Figure 1, within your corpus “computer” and “geek.” Say the first document included four of the word “computer” and one of the word “geek.” Then the second document had two of the word “geek” and one of the word “computer.” You could draw these two very simple documents in two-dimensions. The angle between these documents is a measure of similarity – if the angle is smaller, then the documents are more similar. By taking the cosine of the angle, we normalize the angle to a scale from zero to one.

Of course, the corpuses we will work with will rarely ever have only two words, so we have to imagine the angle between the two documents in much higher dimensions. Cosine




Figure 1:

similarity can be computed simply by standardizing the inner product. First, take each word within each document and divide it by the length of the document. Then, take the inner product of the two resulting vectors.

$$\cos\theta = \left(\frac{Doc1}{||Doc1||} \right) \cdot \left(\frac{Doc2}{||Doc2||} \right)$$

where $||Doc1||$ is the length of the first document. Note that cosine similarity does not change if you simply duplicate the words within one document since the angle between the two vectors will stay the same.

For example, say we are interested in a corpus of President Barack Obama’s press releases from his time in the Senate, from 2005-2008 (Grimmer, 2013). While reading through the documents, we became interested in one particular press release about taxes, entitled “Obama Calls on IRS to Protect Taxpayers Privacy.” We’d like to know whether or not there were more documents in the corpus like this one.

One way to go about this would be to read all of the 709 documents in the corpus. But another way would be to use text as data methods. First, we would remove punctuation, convert to lowercase, remove stopwords, stem, and do whatever other more complicated pre-processing techniques we think might capture the kind of similarities we are looking for. Then we would convert the 709 documents into a document-feature matrix. With a few lines of code, we could calculate the cosine similarity between all other 708 documents and the document we are interested in. We would find that the most similar documents to “Obama Calls on IRS to Protect Taxpayers Privacy” include “Obama Introduces Bill to Stop Tax Preparers from Selling Confidential Tax Information,” “Obama Calls on the TSA to Address Recent Security Vulnerabilities at O Hare,” and “Obama and Coburn Introduce Strengthening Transparency and Accountability in Federal Spending Act of 2008,” all releases that

have to do with privacy and/or taxes.

3.7 Vector representations of words

So far, we have only represented text as a bag of words, or by whether or not they occur within a document. We haven't used any information about the words' relationships with each other or the order in which they appear within the document. When we compute similarities between documents, there are no gradations of how similar words are across documents. For example, if "bus," "car," and "beautiful" are all unique words within the vocabulary, "bus" and "car" are equally similar in the bag of words context as "bus" and "beautiful." But of course, we know much more about any given word – we know the context in which it was used within the document and other words within the language that we are working in that typically appear around it. As we explained before, we could include n-grams within the document-feature matrix to capture some of this context. But in a large corpus, this list of possible n-grams becomes very long and computationally burdensome. Indeed, it may not be important to document every other word a word has ever been used with, but rather to capture something more general about the word itself and its similarities to other words.

3.7.1 Word Embeddings

One way to address this problem is to use word embeddings to encode more information about individual words. The idea behind word embeddings is to project each word in the vocabulary onto an n -dimensional space with all other words. That is, instead of representing words as a binary, 0/1 variable that indicates its presence or absence, instead represent each individual word as a vector with length n . Words that are nearer to each other within this n -dimensional space will be more similar to each other – more likely to be substitutes or synonyms for one another.

Note that this is a much more information-rich way of describing text – instead of describing documents as vectors, we are now describing words, which is a more granular level of analysis – as vectors. The problem with this approach is that typically the unit of analysis of interest to the research is at the *document* level. To address this, the word embeddings can be computed at or aggregated to the document level in order to describe the area of the word space that the document exists in. Documents that are in similar spaces are more likely to be similar to each other. The advantage of this approach is even if the documents do not use the exact same wording, if they use words that are commonly used as substitutes for one another, they will still occupy a similar space and will therefore be calculated as similar.

Word embeddings can be computed in many different ways – from computing word co-occurrence statistics, to computing Principal Components Analysis on the words themselves and Latent Semantic Analysis – but all output a matrix with the number of columns representing each unique word in the vocabulary and the column itself a representation of where the word lies within the space of all vocabulary (Bengio et al., 2006). For purposes of this section, we will go through one increasingly popular method of creating word embeddings, called the skip-gram model (Mikolov et al., 2013). In the skip-gram model, word vectors are computed by setting up a prediction problem where the words surrounding a selected word predict the selected word as if it was missing. For example, "the quick brown fox [blank] over the lazy dog" would use the estimated word vectors of "the, quick, brown, fox, over, the, lazy, dog" to predict the word in the blank. The idea is that words that are used in similar contexts will have similar word vectors; they will predict similar words to fill in the blank. By iterating this window across all words in the corpus, the objective is to maximize the average of the probability of seeing the correct word w_t given the words within the context w_{t-k}, \dots, w_{t+k} , following Le and Mikolov (2014):

$$\frac{1}{T} \sum_{t=k}^{T-k} \log p(w_t | w_{t-k}, \dots, w_{t+k})$$

where $p(w_t | w_{t-k}, \dots, w_{t+k})$ is parameterized as:

$$p(w_t | w_{t-k}, \dots, w_{t+k}) = \frac{e^{y_{w_t}}}{\sum_i e^{y_{w_i}}}$$

where y_{w_i} is the unnormalized probability of seeing any word w_i given the words in the context w_{t-k}, \dots, w_{t+k} , given by

$$y = b + Uh(w_{t-k}, \dots, w_{t+k}; W)$$

Here U and b are coefficients that are optimized, along with W which is a matrix where each column contains the word vector for each unique word. h is a function that averages the word vectors for w_{t-k}, \dots, w_{t+k} that are selected from W . Mikolov et al. (2013) also provide a relatively fast implementation to train word vectors on large corpuses called word2vec.

The output W of word2vec allows us to describe similarities among words within the corpus. We can use cosine similarity as described above to examine similarities between words. For example, we trained a word2vec model on all unigrams and bigrams from a corpus of President Barack Obama’s press releases from his time in the Senate, from 2005-2008. We then looked for the most words closest in the word2vec space to “climate_change” and “pollution.” Figure 2 plots the most similar words in terms of cosine similarity in the word2vec space to these two words. As you can see, some words are used very similar to both climate change and pollution – words such as “emissions” and “greenhouse gas” are highly

similar to both. But then there are words that are more likely to be similar to climate change and not pollution – particularly words about international affairs like “our dependence” and “foreign oil.” Similarly, there are words that are more likely to be used as substitutes for pollution, such as “clean water” or “great lakes” that are less likely to be used as substitutes for climate change. This analysis sheds some insight into the ways in which Obama used the word climate change – in a global context – and alternatively the way he used the word pollution – in a local, Illinois-specific, context.

While word vectors produced by word2vec are typically used as inputs in classification and translation problems, they also have a number of intuitive properties that make them interesting in their own right. Not only do similar words appear in similar places in the space, analogous words appear in parallel planes of the space. This means that arithmetic can be done on the word vectors themselves – in one of the corpuses trained by Mikolov et al. (2013) subtracting the word vector for “man” from the word vector for “king” gives you the address close to the word vector for “queen.”

In the Obama corpus, we can use this arithmetic to understand how Obama used particular words in different contexts during his time in the Senate. For example, the closest words to the vector that results from adding the word vector for “poor” and the word vector for “tax” is “ssi_eligibility,” “gross_income,” and “retirement_plans.” On the other hand, among the closest words to the vector that results from subtracting the word vector for “poor” from the word vector for “tax” is “cheat,” “foreign_subsidary,” “reward_companies,” “an_unfair,” and “bad_actors.” It’s easy to see that the way that Obama talked about taxes and the poor is quite different from the way in which he talked about taxes for the wealthy.

3.7.2 Document Embeddings

While word embeddings retain more information about an individual word, they come at a cost – instead of representing documents, like the document-feature matrix we constructed




Figure 2:

at the beginning of this chapter did, we have now produced representations of words across documents, or a matrix that retains a vector for each unique word within the vocabulary. Because most of the tasks we are interested in in this book are at the document level, we need a way to aggregate the word vector into documents – i.e. to use the information provided by the word embeddings to describe the documents.

There are several ways to aggregate word vectors at the document level. One is to simply take a weighted average of the word vectors that are included within the document, or by some other aggregation process. When this is applied, researchers might use word vectors pre-trained on a very large corpus like Wikipedia to capture word usage in the language as a whole, then aggregate these vectors to describe the documents in their corpus. Or, they can train word vectors within their specific corpus and then aggregate to the document level.

Alternatively, document vectors can be learned explicitly from the data, as in doc2vec (Le and Mikolov, 2014). doc2vec employs a similar model as the skip-gram algorithm, but in addition to the word vectors includes a document-level vector in the prediction model for y . Similarly, the aggregation of words can be an explicit part of a method for classification, as often used in neural networks (described in Chapter 5). Either way, once documents are represented by vectors, we can then perform computations on the document vectors, including computing the similarities between pairs of documents, as discussed before.

3.8 Examples

We now turn to some examples of how the simple creation of the document-feature matrix facilitates analyses that allow us to better understand the content of text and make inferences about the social world. Most simply, as we showed above, taking the column sums of a document-feature matrices allows us to create histograms of the word counts within a corpus. If we then select the N largest column sums, we can quickly understand the most central tendencies of the text. Often, these most frequent words within a corpus are summarized by

words) distinguish the two authors – Hamilton and Madison use words like *upon*, *by*, and *to* at different rates in documents known to be written by the two authors. They decide on a unit of analysis as blocks of text of about 200 words each, in order to be able to distinguish joint authorship within one document. They then count the use of filler words in each of these blocks, similar to our creation of a document-feature matrix.

The authors then create a model that uses the stop words to predict the authorship of papers using existing known documents written by the two authors. They then predict the authorship of the unknown papers using the estimated relationship between stopwords and authorship. The model distinguishes authorship in the unknown documents quite decisively – attributing all 12 documents with unknown authorship to Hamilton. The conclusions of Mosteller and Wallace (1963) have been confirmed in subsequent, more sophisticated analyses of the documents [cites?].

Mosteller and Wallace (1963) are not the only example of authors who have examined a particular type of word in the research process. Jaros and Pan (Forthcoming) use named-entity recognition software in official Chinese newspapers to better understand the distribution of power among political elites in China. Many had speculated that after he took office, General Secretary Xi Jinping consolidated power more than his predecessor, Hu Jintao. Other China watchers observed that the Party had consolidated more power relative to the government in China and that foreign influence seemed to be declining in China.

Of course, all of these observations are difficult to measure. However, in China official newspapers typically reflect Party positions. To measure power, Jaros and Pan (Forthcoming) used named entity recognition in Chinese official newspapers to recognize all people and organization named in Party newspapers. The measure power as the number of times an entity is mentioned, as a proportion of all other entities mentioned. They found – consistent with speculation – that the newspapers reflected many of these trends. Xi Jinping was much more mentioned in Party newspapers than his predecessor was during his time in office.

Party institutions also gained prominence over the time that Xi was in office. However, they show that there was not marketed decline in mentions of foreign organizations.

3.9 What is lost in quantitative analysis? How could this ever work?

When we turn text into numbers, we lose a lot of information. On the one hand, discarding information helps us – it allows us to make the problem of analyzing text as scale tractable. On the other hand, we lose a lot of nuance when we discard information. When we simply count the number of times each document uses “love,” we aren’t sure which of those documents use it genuinely, “I love ice cream,” and which use it facetiously, “We all love homework.” Indeed, these two loves are counted similarly to each other. We also discard the word order – typically more emphasis is on words in the beginning or end of a document, or maybe one of the words was in a heading in 24-point bold font, while another was relegated to a footnote, and these words are still treated equally.

Much about the context of the document we also lose. The words within the text are only that – words – and they may mean something very different in different time periods, social, and political contexts. Infamously, “dog whistles” are identical words or phrases that have two meanings based entirely on the audience, and we may not know the types of people who read the text or what the author intended for them to mean. With political, social, and economic pressure, authors may also be indirect, beat around the bush, or use wordplay or homophones to get their point across in almost “code” and these elements of the text might be lost on the analyst.

So how could this ever work? One answer is that it might not. Throughout this book, we put heavy emphasis on validation – reading the text to make sure they accord with the analysis, re-conducting the analyses in different ways and on different datasets, and exploring

all possible observable implications of the theory or argument. Only after extensive validation can we be sure that we have done any analysis as scientifically as possible.

But also, discarding complexity in the text allows us to organize, compare, and paint a broad picture of texts – which computers are good at – and reserve more times for what we as humans are good at – selecting texts to read and understand thoroughly. The point of much of text analysis methods is to characterize the haystack, to make a broad point about (sometimes) very large corpuses. If one straw of hay doesn’t fit into the broad picture, that usually won’t worry us, as long as we capture the general points in the text more broadly. If we make smart choices in sampling and representing texts quantitatively, we will have more confidence that we have painted the correct birds eye view of our corpus.

Put this in somewhere

References

- Abello, James, Peter Broadwell and Timothy R Tangherlini. 2012. “Computational folkloristics.” *Communications of the ACM* 55(7):60–70.
- Asur, Sitaram and Bernardo A Huberman. 2010. Predicting the future with social media. In *Web Intelligence and Intelligent Agent Technology (WI-IAT), 2010 IEEE/WIC/ACM International Conference on*. Vol. 1 IEEE pp. 492–499.
- Barberá, Pablo and Gonzalo Rivero. 2014. “Understanding the political representativeness of Twitter users.” *Social Science Computer Review* p. 0894439314558836.
- Bengio, Yoshua, Holger Schwenk, Jean-Sébastien Senécal, Frédéric Morin and Jean-Luc Gauvain. 2006. *Neural Probabilistic Language Models*. Berlin, Heidelberg: Springer Berlin Heidelberg pp. 137–186.
- Blaydes, Lisa, Justin Grimmer and Alison McQueen. 2014. Mirrors for princes and sultans:

advice on the art of governance in the medieval Christian and Islamic worlds. Technical report Working Paper.

Boydston, Michelle D. 1991. "Press censorship and access restrictions during the Persian Gulf War: A First Amendment analysis." *Loy. LAL Rev.* 25:1073.

Danzger, M Herbert. 1975. "Validating conflict data." *American Sociological Review* pp. 570–584.

Denny, Matthew James and Arthur Spirling. 2016. "Assessing the consequences of text preprocessing decisions."

Dodds, Peter Sheridan, Kameron Decker Harris, Isabel M Kloumann, Catherine A Bliss and Christopher M Danforth. 2011. "Temporal patterns of happiness and information in a global social network: Hedonometrics and Twitter." *PloS one* 6(12):e26752.

D’Orazio, Vito, Steven T Landis, Glenn Palmer and Philip Schrodtt. 2014. "Separating the wheat from the chaff: Applications of automated document classification using support vector machines." *Political analysis* 22(2):224–242.

Franco, Annie, Justin Grimmer and Monica Lee. 2016. "Changing the Subject to Build an Audience: How Elected Officials Affect Constituent Communication."

Gill, Michael and Arthur Spirling. 2015. "Estimating the Severity of the WikiLeaks US Diplomatic Cables Disclosure." *Political Analysis* 23(2):299–305.

Golder, Scott A and Michael W Macy. 2011. "Diurnal and seasonal mood vary with work, sleep, and daylength across diverse cultures." *Science* 333(6051):1878–1881.

Graesser, Arthur C., Danielle S. McNamara, Max M. Louwerse and Zhiqiang Cai. 2004. "Coh-Metrix: Analysis of text on cohesion and language." *Behavior Research Methods*,

Instruments, & Computers 36(2):193–202.

URL: <https://doi.org/10.3758/BF03195564>

Grimmer, Justin. 2013. *Representational Style in Congress: What Legislators Say and Why It Matters*. Cambridge University Press.

Grimmer, Justin and Brandon M Stewart. 2013. “Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts.” *Political analysis* 21(3):267–297.

Gup, Ted. 2008. *Nation of Secrets: The Threat to Democracy and the American Way of Life*. Anchor.

Handler, Abram, Matthew J Denny, Hanna Wallach and Brendan O’Connor. 2016. “Bag of what? simple noun phrase extraction for text analysis.” *NLP+ CSS* 114.

House, Freedom. 2015. *Freedom on the Net 2015:Privatising Censorship, Eroding Privacy*.

Jaros, Kyle and Jennifer Pan. Forthcoming. “China’s Newsmakers: Official Media Coverage and Political Shifts in the Xi Jinping Era.” *China Quarterly* .

Jursfsky, Dan and James Martin. 2009. *Speech and natural language processing: An introduction to natural language processing, computational linguistics, and speech recognition*. Upper Saddle River: Prentice Hall.

King, Gary, Jennifer Pan and Margaret E. Roberts. 2013. “How Censorship in China Allows Government Criticism but Silences Collective Expression.” *American Political Science Review* 107:1–18. <http://j.mp/LdVXqN>.

King, Gary, Patrick Lam and Margaret E Roberts. 2014. “Computer-Assisted Keyword and Document Set Discovery from Unstructured Text.”.

- Lampos, Vasileios and Nello Cristianini. 2010. Tracking the flu pandemic by monitoring the social web. In *2010 2nd International Workshop on Cognitive Information Processing*. IEEE pp. 411–416.
- Le, Quoc and Tomas Mikolov. 2014. Distributed Representations of Sentences and Documents. In *Proceedings of the 31st International Conference on Machine Learning*, ed. Eric P. Xing and Tony Jebara. Vol. 32 of *Proceedings of Machine Learning Research* Beijing, China: PMLR pp. 1188–1196.
URL: <http://proceedings.mlr.press/v32/le14.html>
- MacKinnon, Rebecca. 2008. “Flatter world and thicker walls? Blogs, censorship and civic discourse in China.” *Public Choice* 134(1-2):31–46.
- Malesky, Edmund, Paul Schuler and Anh Tran. 2012. “The adverse effects of sunshine: a field experiment on legislative transparency in an authoritarian assembly.” *American Political Science Review* 106(04):762–786.
- Manning, Christopher D., Mihai Surdeanu, John Bauer, Jenny Finkel, Steven J. Bethard and David McClosky. 2014. The Stanford CoreNLP Natural Language Processing Toolkit. In *Association for Computational Linguistics (ACL) System Demonstrations*. pp. 55–60.
URL: <http://www.aclweb.org/anthology/P/P14/P14-5010>
- Manning, Christopher D., Prabhakar Raghavan and Hinrich Schütze. 2008. *Introduction to Information Retrieval*. NY: Cambridge University Press.
- Marcus, Mitchell P., Mary Ann Marcinkiewicz and Beatrice Santorini. 1993. “Building a Large Annotated Corpus of English: The Penn Treebank.” *Comput. Linguist.* 19(2):313–330.
URL: <http://dl.acm.org/citation.cfm?id=972470.972475>

- Mikolov, Tomas, Ilya Sutskever, Kai Chen, Greg S Corrado and Jeff Dean. 2013. Distributed Representations of Words and Phrases and their Compositionality. In *Advances in Neural Information Processing Systems 26*, ed. C. J. C. Burges, L. Bottou, M. Welling, Z. Ghahramani and K. Q. Weinberger. Curran Associates, Inc. pp. 3111–3119.
- URL:** <http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>
- Miller, George A., Richard Beckwith, Christiane Fellbaum, Derek Gross and Katherine J. Miller. 1990. “Introduction to WordNet: An On-line Lexical Database*.” *International Journal of Lexicography* 3(4):235–244.
- URL:** + <http://dx.doi.org/10.1093/ijl/3.4.235>
- Mosteller, Frederick and David L Wallace. 1963. “Inference in an authorship problem: A comparative study of discrimination methods applied to the authorship of the disputed Federalist Papers.” *Journal of the American Statistical Association* 58(302):275–309.
- Norris, Pippa. 2001. *Digital divide: Civic engagement, information poverty, and the Internet worldwide*. Cambridge University Press.
- O’Connor, Brendan, Brandon M. Stewart and Noah A. Smith. 2013. Learning to Extract International Relations from Political Context. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Sofia, Bulgaria: Association for Computational Linguistics pp. 1094–1104.
- URL:** <http://www.aclweb.org/anthology/P13-1108>
- O’Connor, Brendan, Ramnath Balasubramanyan, Bryan R Routledge and Noah A Smith. 2010. “From tweets to polls: Linking text sentiment to public opinion time series.” *ICWSM* 11(122-129):1–2.

- Pomerantsev, P. 2014. *Nothing Is True and Everything Is Possible: The Surreal Heart of the New Russia*. PublicAffairs.
- Porter, Martin F. 1980. “An algorithm for suffix stripping.” *Program* 14(3):130–137.
- Schonhardt-Bailey, Cheryl. N.d. *Accountability, Oversight and Deliberation of Economic Policy in UK Parliamentary Committees*.
- Snyder, David and William R Kelly. 1977. “Conflict intensity, media sensitivity and the validity of newspaper data.” *American Sociological Review* pp. 105–123.
- Strapparava, Carlo and Rada Mihalcea. 2008. Learning to Identify Emotions in Text. In *Proceedings of the 2008 ACM Symposium on Applied Computing*. SAC '08 New York, NY, USA: ACM pp. 1556–1560.
- URL:** <http://doi.acm.org/10.1145/1363686.1364052>